

ropls: PCA, PLS(-DA) and OPLS(-DA) for multivariate analysis and feature selection of omics data

Etienne A. Thévenot

August 7, 2015

1 The *ropls* package

The *ropls* R package implements the PCA, PLS(-DA) and OPLS(-DA) approaches with the original, NIPALS-based, versions of the algorithms [1][2]. It includes the R² and Q² quality metrics [3][4], the permutation diagnostics [5], the computation of the VIP values [6], the score and orthogonal distances to detect outliers [7], as well as many graphics (scores, loadings, predictions, diagnostics, outliers, etc).

The functionalities from *ropls* can also be accessed via a graphical user interface in the *Multivariate* module from the *Workflow4Metabolomics.org* online resource for computational metabolomics (<http://workflow4metabolomics.org/>) which provides a user-friendly, web-based environment for data pre-processing, statistical analysis, and annotation [8].

2 Context

2.1 Orthogonal Partial Least-Squares

Partial least-squares, which is a latent variable regression method based on covariance between the predictors and the response, has been shown to efficiently handle datasets with multi-collinear predictors, as in the case of spectrometry measurements [1]. More recently, Trygg and Wold introduced the orthogonal projection to latent structures (OPLS) algorithm to model separately the variations of the predictors correlated and orthogonal to the response [2].

OPLS has a similar predictive capacity compared to PLS and improves the interpretation of the predictive components and of the systematic variation [9]. In particular, OPLS modeling of single responses only requires one predictive component.

Diagnostics such as the Q²Y metrics and permutation testing are of high importance to avoid overfitting and assess the statistical significance of the model. The variable importance in projection (VIP), which reflects both the loading weights for each component and the variability of the response explained by this component [9][10] is often used for feature selection [2][9].

2.2 OPLS software

OPLS is available in the SIMCA-P commercial software (Umetrics, Umeå, Sweden; [3]). In addition, the kernel-based version of OPLS [11] is available in the open-source R statistical environment [12], and a single implementation of the linear algorithm in R has been described recently [13].

3 The 'sacurine' metabolomics dataset

3.1 Study objective

The objective was to study the influence of age, body mass index and gender on metabolite concentrations in urine, by analysing 183 samples from a cohort of adults with liquid chromatography coupled to high-resolution mass spectrometry ([14]).

3.2 Pre-processing and annotation

Urine samples were analyzed by using an LTQ Orbitrap in the negative ionization mode. A total of 109 metabolites were identified or annotated at the MSI level 1 or 2. After retention time alignment with XCMS, peaks were integrated with Quan Browser. Signal drift and batch effect were corrected, and each urine profile was normalized to the osmolality of the sample. Finally, the data were log10 transformed.

3.3 Covariates

The volunteers' age, body mass index, and gender were recorded.

4 Hands-on

4.1 Loading

We first load the *ropls* package:

```
> library(ropls)
```

We then load the *sacurine* dataset which contains:

1. The *dataMatrix* matrix of numeric type containing the intensity profiles (log10 transformed)
2. The *sampleMetadata* data frame containing sample metadata
3. The *variableMetadata* data frame containing variable metadata

```
> data(sacurine)
```

```
> names(sacurine)
```

```
[1] "dataMatrix"      "sampleMetadata"  "variableMetadata"
```

We attach *sacurine* to the search path and display a summary of the content of the *dataMatrix*, *sampleMetadata* and *variableMetadata* with the *strF* Function of the *ropls* package (see also *str*)

```
> attach(sacurine)
```

```
> strF(dataMatrix)
```

```

      dim class   mode typeof      size NAs  min mean median max
183 x 109 matrix numeric double 178800 bytes  0 -0.3  4.2   4.3   6
(2-methoxyethoxy)propanoic acid isomer (gamma)Glu-Leu/Ile ...
HU_011                3.019766011          3.888479324 ...
HU_014                3.81433889          4.277148905 ...
...
HU_208                3.748127215          4.523763202 ...
HU_209                4.208859398          4.675880567 ...
Valerylglycine isomer 2 Xanthosine
HU_011                3.889078716 4.075879575
HU_014                4.181765852 4.195761901
...
HU_208                4.634338821 4.487781609
HU_209                4.47194762 4.222953354

> strF(sampleMetadata)

  ageVn  bmiVn ...  bmiVn genderFc
numeric numeric ... numeric  factor
nRow nCol      size NAs
183   3 15280 bytes  0
  ageVn bmiVn genderFc
HU_011   29 19.75      M
HU_014   59 22.64      F
...      ...      ...  <NA>
HU_208   27 18.61      F
HU_209  17.5 21.48      F

> strF(variableMetadata)

msiLevelVn  hmdbVc ...  hmdbVc chemicalClassVc
  numeric character ... character      character
nRow nCol      size NAs
109   3 16352 bytes  0

      msiLevelVn  hmdbVc chemicalClassVc
(2-methoxyethoxy)propanoic acid isomer      2      Organi
(gamma)Glu-Leu/Ile      2      AA-pep
...      ...      ...
Valerylglycine isomer 2      2      AA-pep:AcyGly
Xanthosine      1 HMDB00299      Nucleo

```

4.2 Principal Component Analysis (PCA)

We perform a PCA on the *dataMatrix* matrix (samples as rows, variables as columns):

```

> sacurine.pca <- opls(dataMatrix)

R2X(cum) pre ort
h8    0.501  8    0

```

A summary of the model (8 components were selected) is printed. In addition a 4-plot figure displays an inertia barplot, outlier diagnostics, and the score and loading plots. To plot the scores only (Figure 1):

```

> plot(sacurine.pca, plotVc = "x-score")

```




Figure 1: PCA score plot. The displayed components can be specified with `parCompVi` (plotting **parameter** specifying the **Components**: **Vector** of 2 integers)

Note:

1. Since `dataMatrix` does not contain missing value, the singular value decomposition was used by default; NIPALS can be selected with the `algoC` argument specifying the **algorithm** (**Character**),
2. The `predI = NA` default number of **predictive** components (**Integer**) means that only components with a variance superior to the mean variance of all components are kept (note that this rule requires all components to be computed and can be quite time-consuming for large datasets with the NIPALS algorithm; in such cases, one may specify a limited number of components with the `predI` parameter).

Let us see if we notice any partition according to gender, by labeling/coloring the samples according to the gender and drawing the Mahalanobis ellipses for the male and female subgroups (Figure 2).

```
> genderFc <- sampleMetadata[, "genderFc"]
> plot(sacurine.pca, plotVc = "x-score",
+ parAsColFcVn = genderFc, parEllipsesL = TRUE)
```

Note that the plotting **parameter** to be used **As Colors** (**Factor** of **character** type or **Vector** of **numeric** type) has a length equal to the number of rows of the `dataMatrix` matrix (ie of samples) and that this qualitative or quantitative variable is converted into colors (by using an internal palette or color scale, respectively). We could have visualized the age of the individuals by specifying `parAsColFcVn = sampleMetadata[, "ageVn"]`.

4.3 Partial least-squares: PLS and PLS-DA

For PLS (and OPLS), the Y response(s) must be provided. Y can be either a numeric vector (respectively matrix) for single (respectively multiple) (O)PLS regression, or a character factor for (O)PLS-DA classification as in the following example (The score plot is displayed on Figure 3).




Figure 2: PCA score plot colored according to gender.

```
> sacurine.plsda <- opls(dataMatrix, genderFc)
```

```
      R2X(cum) R2Y(cum) Q2(cum) RMSEE pre ort
h3      0.275      0.73  0.584 0.262  3  0
```

Note:

1. The NIPALS algorithm is used for PLS (and OPLS); *dataMatrix* matrices with (a moderate amount of) missing values can thus be analysed,
2. As for PCA, the default *predI* value is NA: components are computed until either R2Y of the component is < 0.01 or Q2Y is < 0 (for more than 100 observations) or 0.05 otherwise.

We see that 3 predictive ('pre') components were computed and that the (cum)ulative R2Y and Q2Y are quite high. Let us assess the significance of these values by comparing to models generated after random permutation of the labels (Figure 4).

```
> sacurine.plsPer <- opls(dataMatrix, genderFc,
+ permI = 100)
```

The R2Y and Q2Y values of the model trained on the models with randomly permuted response labels are all below the values of the model with the true response (horizontal lines). Significance of our PLS-DA model should be further confirmed by using a higher number of permutations.

4.4 Orthogonal partial least squares: OPLS and OPLS-DA

To perform OPLS(-DA), we set *orthoI* (number of components which are **orthogonal**; Integer) to either a specific number of orthogonal components, or to NA. Let us build an OPLS-DA model of the 'gender' response (Figure 4). Note that for OPLS modeling of a single response, the number of predictive component is 1.




Figure 3: PLS-DA model of the 'gender' response (score plot). The number of components and the cumulative R2X, R2Y and Q2Y are indicated below the plot.




Figure 4: Significance of R2Y and Q2Y evaluated by permutation testing.

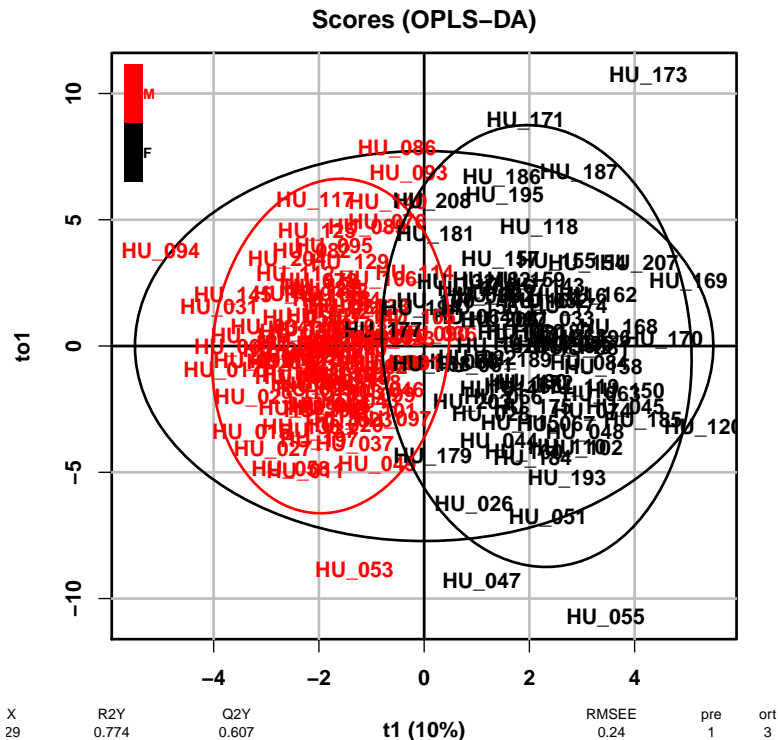


Figure 5: OPLS-DA model of the 'gender' response (score plot). The predictive component is displayed as abscissa and the orthogonal component as ordinate.

```
> sacurine.oplsda <- opls(dataMatrix, genderFc,
+ predI = 1, orthoI = NA)
```

	R2X(cum)	R2Y(cum)	Q2(cum)	RMSEE	pre	ort
sum	0.329	0.774	0.607	0.24	1	3

Let us assess the predictive performance of our model. We first train the model on a subset of the samples (here we use the 'odd' subset value which splits the data set into two halves with similar proportions of samples for each class; alternatively, we could have used a specific subset of indices for training):

```
> sacurine.oplsda <- opls(dataMatrix, genderFc, predI = 1, orthoI = NA, subset = "odd")
```

	R2X(cum)	R2Y(cum)	Q2(cum)	RMSEE	RMSEP	pre	ort
sum	0.18	0.767	0.563	0.245	0.342	1	1

We compute the performances on the training subset:

```
> trainVi <- sacurine.oplsda[["subset"]]
> table(genderFc[trainVi], fitted(sacurine.oplsda))
```

	M	F
M	49	1
F	3	39

We then compute the performances on the test subset:

```
> table(genderFc[-trainVi],
+ predict(sacurine.oplsda, dataMatrix[-trainVi, ]))
```

	M	F
M	46	4

F 4 37

As expected, the predictions on the test subset are slightly lower. The classifier however still achieves 91% of correct predictions.

Before closing this example session, we detach *sacurine* from the search path:

```
> detach(sacurine)
```

4.5 Session information

```
> sessionInfo()
```

- R version 3.2.0 (2015-04-16), x86_64-w64-mingw32
- Locale: LC_COLLATE=French_France.1252, LC_CTYPE=French_France.1252, LC_MONETARY=French_France.1252, LC_NUMERIC=C, LC_TIME=French_France.1252
- Base packages: base, datasets, graphics, grDevices, methods, stats, utils
- Other packages: ropls 1.1.3
- Loaded via a namespace (and not attached): BiocStyle 1.6.0, tools 3.2.0

5 Pre-processing and annotation of mass spectrometry data

To illustrate how *dataMatrix*, *sampleMetadata* and *variableMetadata* can be obtained from raw mass spectra file, we use the LC-MS data from the *faahKO* package [15], which will pre-process with the *xcms* package [16] and annotate isotopes and adducts with the *CAMERA* package [17], as described in the corresponding vignettes (all these packages are from *bioconductor*).

Let us start by getting the paths to the 12 raw files (6 KO and 6 WT mice) in the ".cdf" open format. The files are grouped in two sub-directories ("KO" and "WT") since *xcms* can use sample class information when grouping the peaks and correcting retention times.

```
> library(faahKO)
> cdfpath <- system.file("cdf", package = "faahKO")
> cdffiles <- list.files(cdfpath, recursive = TRUE, full.names = TRUE)
> basename(cdffiles)
```

Next, *xcms* is used to pre-process the individual raw files, as described in the vignette.

```
> library(xcms)
> xset <- xcmsSet(cdffiles)
> xset
> xset <- group(xset)
> xset2 <- retcor(xset, family = "symmetric", plottype = "mdevden")
> xset2 <- group(xset2, bw = 10)
> xset3 <- fillPeaks(xset2)
```

Finally, the *annotateDiffreport* from *CAMERA* annotates isotopes and adducts and builds a peak table containing the peak intensities and the variable metadata.

```
> library(CAMERA)
> diffreport <- annotateDiffreport(xset3, quick=TRUE)
> diffreport[1:4, ]
```

We then build the *dataMatrix*, *sampleMetadata* and *variableMetadata* matrix and dataframes as follows:


```
> sampleVc <- grep("^ko|^wt", colnames(diffreport), value = TRUE)
> dataMatrix <- t(as.matrix(diffreport[, sampleVc]))
> dimnames(dataMatrix) <- list(sampleVc, diffreport[, "name"])
> sampleMetadata <- data.frame(row.names = sampleVc,
+ genotypeFc = substr(sampleVc, 1, 2))
> variableMetadata <- diffreport[, !(colnames(diffreport) %in% c("name", sampleVc))]
> rownames(variableMetadata) <- diffreport[, "name"]
```

The data can now be analysed with the *ropIs* package as described in the previous section (i.e. by performing a PCA and a PLS-DA):

```
> library(ropIs)
> opIs(dataMatrix)
> opIs(dataMatrix, sampleMetadata[, "genotypeFc"])
```

6 Other datasets

In addition to the *sacurine* dataset presented above, the package contains the following datasets to illustrate the functionalities of PCA, PLS and OPLS (see the examples in the documentation of the *opIs* function):

aminoacids Amino-Acids Dataset. Quantitative structure property relationship (QSPR; [1]).

cellulose NIR-Viscosity example data set to illustrate multivariate calibration using PLS, spectral filtering and OPLS (Multivariate calibration using spectral data. Simca tutorial. Umetrics, Sweden).

cornell Octane of various blends of gasoline: Twelve mixture component proportions of the blend are analysed [4].

foods Food consumption patterns accross European countries (FOODS). The relative consumption of 20 food items was compiled for 16 countries. The values range between 0 and 100 percent and a high value corresponds to a high consumption. The dataset contains 3 missing data [3].

linnerud Three physiological and three exercise variables are measured on twenty middle-aged men in a fitness club [4].

lowarp A multi response optimization data set (LOWARP) [3].

mark Marks obtained by french students in mathematics, physics, french and english. Toy example to illustrate the potentialities of PCA [18].

References

- [1] S. Wold, M. Sjöström, and L. Eriksson. Pls-regression: a basic tool of chemometrics. *Chemometrics and Intelligent Laboratory Systems*, 58:109–130, 2001. URL: [http://dx.doi.org/10.1016/S0169-7439\(01\)00155-1](http://dx.doi.org/10.1016/S0169-7439(01)00155-1).

- [2] J. Trygg and S. Wold. Orthogonal projection to latent structures (o-pls). *Journal of Chemometrics*, 16:119–128, 2002. URL: <http://dx.doi.org/10.1002/cem.695>.
- [3] I. Eriksson, E. Johansson, N. Kettaneh-Wold, and S. Wold. *Multi- and megavariate data analysis. Principles and applications*. Umetrics Academy, 2001.
- [4] M. Tenenhaus. *La regression PLS : theorie et pratique*. Editions Technip, 1998.
- [5] Ewa Szymanska, Edoardo Saccenti, AgeK. Smilde, and JohanA. Westerhuis. Double-check: validation of diagnostic statistics for pls-da models in metabolomics studies. *Metabolomics*, 8:3–16, 2012. URL: <http://dx.doi.org/10.1007/s11306-011-0330-3>, doi:10.1007/s11306-011-0330-3.
- [6] B.-H. Mevik and R. Wehrens. The pls package: Principal component and partial least squares regression in r. *Journal of Statistical Software*, 18, 2007.
- [7] M. Hubert, P.J. Rousseeuw, and K. Vanden Branden. Robpca: a new approach to robust principal component analysis. *Technometrics*, 47:64–79, 2005.
- [8] Franck Giacomoni, Gildas Le Corguillé, Misharl Monsoor, Marion Landi, Pierre Pericard, Mélanie Pétéra, Christophe Duperier, Marie Tremblay-Franco, Jean-François Martin, Daniel Jacob, Sophie Goulitquer, Etienne A. Thévenot, and Christophe Caron. Workflow4metabolomics: a collaborative research infrastructure for computational metabolomics. *Bioinformatics*, 31(9):1493–1495, 2015. URL: <http://bioinformatics.oxfordjournals.org/content/31/9/1493.abstract>, arXiv: <http://bioinformatics.oxfordjournals.org/content/31/9/1493.full.pdf+html>.
- [9] Rui Climaco Pinto, Johan Trygg, and Johan Gottfries. Advantages of orthogonal inspection in chemometrics. *Journal of Chemometrics*, 26(6):231–235, 2012. URL: <http://dx.doi.org/10.1002/cem.2441>, doi:10.1002/cem.2441.
- [10] Tahir Mehmood, Kristian Hovde Liland, Lars Snipen, and Solve Saebo. A review of variable selection methods in partial least squares regression. *Chemometrics and Intelligent Laboratory Systems*, 118(0):62–69, 2012. URL: <http://www.sciencedirect.com/science/article/pii/S0169743912001542>, doi:http://dx.doi.org/10.1016/j.chemolab.2012.07.010.
- [11] Max Bylesjo, Mattias Rantalainen, Jeremy Nicholson, Elaine Holmes, and Johan Trygg. K-opls package: Kernel-based orthogonal projections to latent structures for prediction and interpretation in feature space. *BMC Bioinformatics*, 9(1):106, 2008. URL: <http://www.biomedcentral.com/1471-2105/9/106>, doi:10.1186/1471-2105-9-106.
- [12] R Development Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2008. ISBN 3-900051-07-0. URL: <http://www.R-project.org>.
- [13] R. Gaude, F. Chignola, D. Spiliotopoulos, A. Spitaleri, M. Ghitti, JM. Garcia-Manteiga, S. Mari, and G. Musco. muma, an r package for metabolomics univariate and multivariate statistical analysis. *Current Metabolomics*, 1:180–189, 2013.
- [14] Etienne A. Thévenot, Aurélie Roux, Xu Ying, Eric Ezan, and Christophe Junot. Analysis of the human adult urinary metabolome variations with age, body mass index and gender by implementing a comprehensive workflow for univariate and opls statistical analyses. *Journal of Proteome Research*, 14(8):3322–3335, 2015. doi:10.1021/acs.jproteome.5b00354.
- [15] Alan Saghatelian, Sunia A. Trauger, Elizabeth J. Want, Edward G. Hawkins, Gary Siuzdak, and Benjamin F. Cravatt. Assignment of endogenous substrates to enzymes by global metabolite profiling. *Biochemistry*, 43(45):14332–14339, November 2004. URL: <http://dx.doi.org/10.1021/bi0480335>.
- [16] C. A. Smith, E. J. Want, G. O’Maille, R. Abagyan, and G. Siuzdak. Xcms: Processing mass spectrometry data for metabolite profiling using nonlinear peak alignment, matching, and identifica-

tion. *Analytical Chemistry*, 78(3):779–787, 2006. URL: <http://dx.doi.org/10.1021/ac051437y>, doi:10.1021/ac051437y.

- [17] Carsten Kuhl, Ralf Tautenhahn, Christoph Bottcher, Tony R. Larson, and Steffen Neumann. Camera: An integrated strategy for compound spectra extraction and annotation of liquid chromatography/mass spectrometry data sets. *Analytical Chemistry*, 84(1):283–289, 2012. URL: <http://pubs.acs.org/doi/abs/10.1021/ac202450g>, arXiv:<http://pubs.acs.org/doi/pdf/10.1021/ac202450g>, doi:10.1021/ac202450g.
- [18] A. Baccini. Statistique descriptive multidimensionnelle (pour les nuls), 2010.